

**Question 1: (3 mark)**

Arrange in ascending order as per the increasing runtime management cost of these entities: (One line justification for each entity)

**Processes, ULTs, Tasks, Threads, and Coroutines**

**Answer:**

Tasks → Coroutines → ULTs → Threads → Processes //+0.5 marks (only if all entities are arranged correctly)

Each correct justification below carrier 0.5 marks (5 x 0.5)

- Task is only a combination of function pointer and arguments to this function.
- Coroutines preserve its current execution state and support context switching, but without an entire thread stack (stack required only for the current function).
- ULTs have all the states associated with a thread, but are managed inside the user space.
- Threads inside a process shares parent process's address space, but are managed by the kernel.
- Processes do not share address space with each other, and are managed by the kernel.

**Question 2: (1 marks)**

Briefly justify if concurrency is the same as parallelism.

**Answer:**

Concurrency is not same as parallelism. //+0.5 marks

Concurrency is ability to divide a program execution into small chunks where these chunks supports interleaved execution. //+0.25 marks

Parallelism is when the execution of above chunks are actually interleaved, i.e., being executed in parallel. //+0.25 marks

**Question 3: (2 marks)**

Give two reasons: Why the cost of context switches performed by the OS is NOT insignificant with large number of running processes/threads.

**Answer:**

- Context switch requires switching the execution from user space (user stack) into kernel space (kernel stack), and it happens frequently with increasing number of threads/processes. //+1 marks
- Frequent invocation of kernel scheduler to decide which thread/process to run next. //+1 marks

**Question 4: (1 marks)**

Which is more scalable for launching 4 tasks on a 4-core processor: work-sharing or work-stealing? (No justification required) **Answer:**

Both will behave the same (+0.5 marks) as creating four tasks will not have any observable overheads unlike creating several tasks. (+0.5 marks)

**Question 5: (3 marks)**

What will be the output of the following program? (No justification required).

```
std::mutex mt;
int main() {
    boost::fibers::fiber F1 ([=]() {
        cout<< "Stage-1\n";
        boost::this_fiber::yield();
        int count=1;
        boost::fibers::fiber F2 ([&]() {
            cout<< "Stage-2\n";
            std::unique_lock<std::mutex> lk(mt);
            count+=2;
        });
        cout<< "Stage-3\n";
        std::unique_lock<std::mutex> lk(mt);
        count+=1;
        F2.join();
        cout<<"count="<<count<<"\n";
    });
    boost::fibers::fiber F3([=]() {
        cout<< "Stage-4\n";
    });
    F3.join();
    F1.join();
    cout<< "Stage-5\n";
}
```

**Answer:**

Stage-1  
Stage-4  
Stage-3  
Stage-2  
<HANG> //as F2 is trying to take a lock held by F1

Each line from Stage-1 to Stage-2 carries +0.5 marks.  
Notifying a hang at correct place carries +1 marks.